

Bangladesh University of Engineering and Technology

CSE 306

Computer Architecture Sessional

Offline 2: 32-bit Floating Point Adder Circuit

Section - B1

Group - 5

2005061 - Farriha Afnan

2005072 - Abrar Rahman Abir

2005079 - Ananya Shahrin Promi

2005087 - Iffat Bin Hossain

2005089 - Wahid Al Azad Navid

January 14, 2024

1 Introduction

Adding floating-point numbers is a big deal in scientific calculations. It is tricky because we need to be precise with our numbers.

A floating-point number has three parts: a sign (plus or minus), an exponent (like a power), and a fraction (some bits after the dot). These parts come together using a special formula to create the floating-point number. This formula ensures that our number is in a standard form, and the exponent is always positive.

$$(-1)^{\text{sign}} \cdot (1 + \text{Fraction}) \cdot 2^{\text{Exponent} - \text{Bias}}$$

When we use a tool called a floating-point adder, we put in two floating-point numbers and get a sum as the result. But when we're adding, we have to be careful about the form of our numbers and make sure we're not losing any important details.

2 Problem Specification

In this assignment, we are required to design a floating point adder circuit that takes two floating points as inputs and provides their sum, another floating point as output. Each floating point will be 32 bits long with the following representation:

Sign	Exponent	Fraction
1 Bit	11 Bit	20 Bit

Table 1: Problem Specification

Input is in IEEE 754 Standard format of binary representation of the floating point number (sign bit, exponent bit, fraction bits). The number obtained after addition is required to be in normal form. Besides, we need to ensure that the number is rounded off properly.

3 Flowchart of the Addition/Subtraction Algorithm

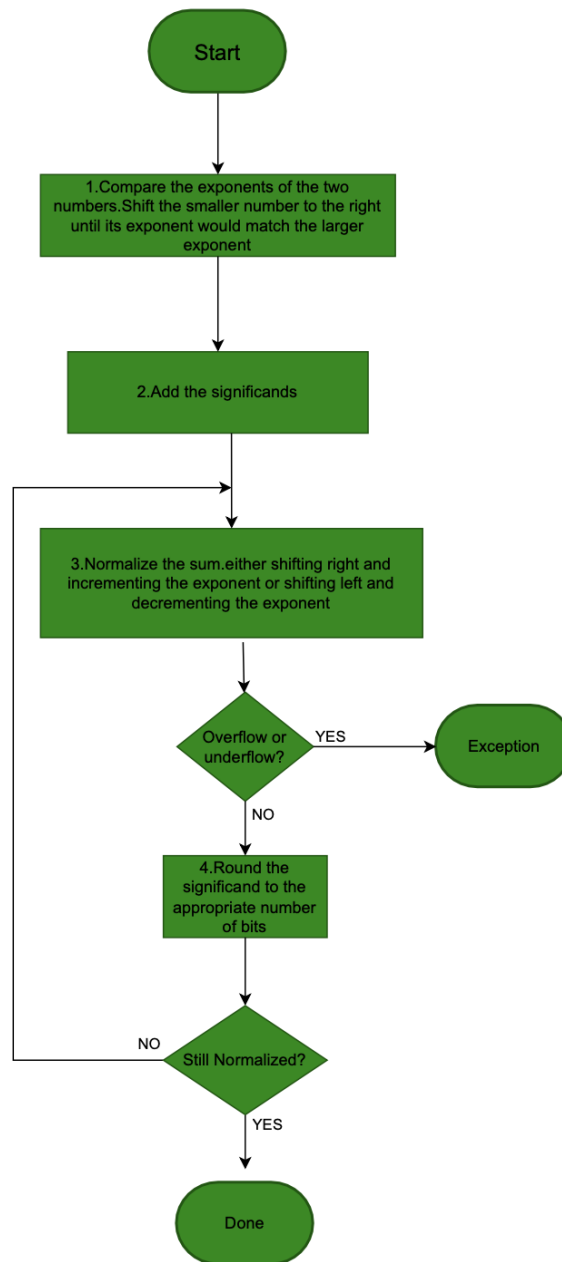


Figure 1: Flowchart of the Algorithm

4 High-Level Block Diagram of the Architecture

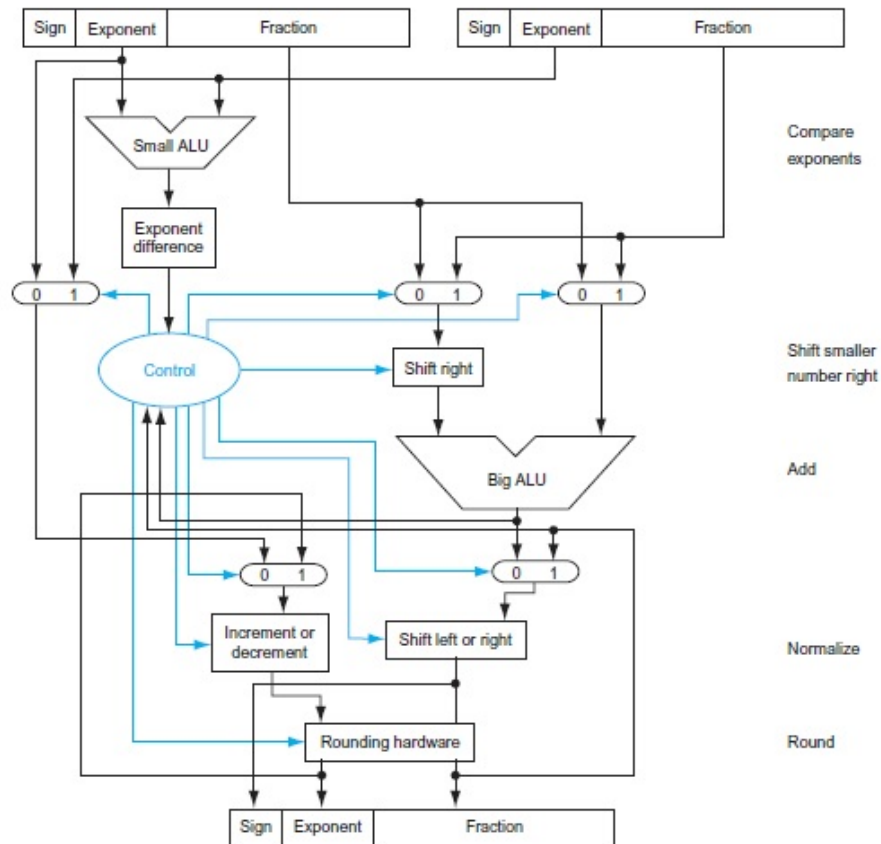


Figure 2: High-Level Block Diagram

5 Detailed Circuit Diagram of the Important Blocks

We will explain the design steps by breaking down and detailing each important component.

5.1 Multiplexer Library

The n -bit 2×1 MUX circuits take two n -bit inputs and output the one selected by a one-bit selector, constructed using cascading IC74157.

5.2 AND Library

The n -Bit $m \times 1$ AND circuits take m n -bit inputs and output their bitwise AND. These are constructed using cascading IC7408.

5.3 OR Library

The n -Bit $m \times 1$ OR circuits take m n -bit inputs and output their bitwise OR. These are constructed using cascading IC7432.

5.4 Exclusive OR Library

n-bit $m \times 1$ XOR takes m n-bit inputs and outputs their bitwise XOR. These are constructed using cascading IC7486.

5.5 NOT Library

n-bit NOT circuits take n-bit input and output its bitwise NOT. These are constructed using cascading IC7404.

5.6 Adder Library

n-bit Adder takes two n-bit inputs and a carry-in bit and outputs their bitwise sum and the carry-out bit, constructed using cascading IC7483.

5.7 Value Add to LSB

It analyzes the 4 least significant bits of a number and outputs the value that needs to be added to the new LSB.




Figure 3: AdderBit

5.8 Compare Input with 4

This circuit takes a 3-bit input and outputs whether the input is greater, equal, or less than 4.




Figure 4: CompareWith4

5.9 Input Splitter

It splits incoming 32-bit data into 3 outgoing wires. First, two 32-bit inputs, A and B are inserted into InputSplitter. This will split the 32-bits into three individual lines of 1-bit (sign bit A_s and B_s), 11-bit (exponent A_e and B_e), and 20-bit (fraction part A_f and B_f) each.

which adds a leading one at the leftmost bit of each fraction and converts the 21-bit fractions into 32-bit values.




Figure 7: MakeSignificand

5.12 Left Shifter

It takes a 32-bit input and an 11-bit input (shift amount). CustomLeftShifter outputs the input left shifted by that amount, with the shifted bits replaced by the shift bit.

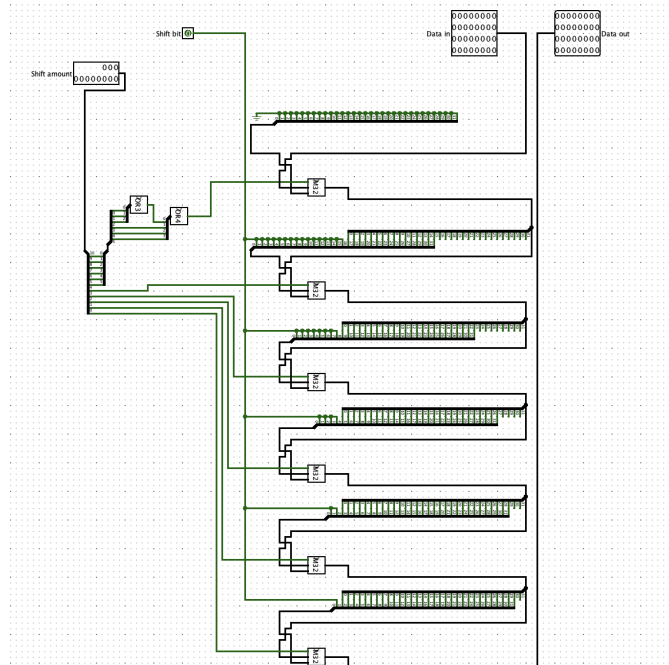


Figure 8: CustomLeftShifter

5.13 Right Shifter

This circuit takes a 32-bit input and an 11-bit input (shift amount) and outputs the input right shifted by that amount, with the shifted bits replaced by the shift bit. Thus, inserting L_f and D_e will shift L_f by an amount of D_e , making $G_e = L_e$.

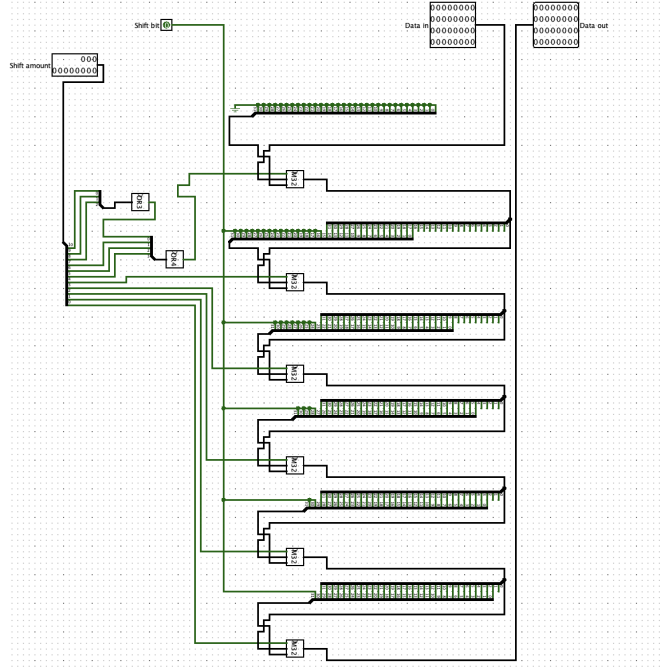


Figure 9: CustomRightShifter

5.14 Inverter

It takes a 32-bit input and a selector and outputs the same input if the selector is set to 0, otherwise outputs the bitwise not (1's complement) of the input.

In subtraction operations, regardless of the signs, we consistently subtract L from G . This




Figure 10: SelectInverse

approach eliminates the need to invert the result after addition, and the sign of the result remains G_s . To achieve this, we utilized the SelectInverse circuit, which outputs the 2's complement of the right-shifted L_f only when $G_e \oplus L_e$ is 1; otherwise, it outputs L_f as is. Moreover, C_{in} should be 1 during the subtraction operation.

5.15 Increment Exponent

We use IncrementExponent to achieve this $F_e = G_e + 1$.




Figure 11: IncrementExponent

5.16 Priority Encoder

PriorityEncoder takes a 32-bit input outputs the position of the first 1 in the input and adds one with this.




Figure 12: PriorityEncoder

5.17 Normalizer

When there is a carry-out, indicating that the leftmost 1 is already beyond the result fraction (we denote it as F_{f32}), and thus, no further right shifting is needed. In the absence of a carry-out, we must left-shift the result until the leftmost 1 is outside the 32 bits. Regardless of whether C_{out} of A_{32} is 1, we transmit F_{f32} and F_e into the Normalizer. Here, using the PriorityEncoder, we ascertain the required shift amount D_{norm} . We then pass D_{norm} and F_{f32} to the CustomLeftShifter and subtract D_{norm} from F_e . Using a pair of MUX, we decide based on C_{out} of A_{32} whether to output the original value or the shifted one. Thus, we obtain the final normalized, unrounded $1 + 11 + 32 = 44$ -bit result. We later checked whether to use the normalized output or not. This was determined according to the truth table below: (Normalized output taken when flag = 0)

G_s	L_s	C_{out}	flag
0	0	0	0
0	0	1	1
0	1	0	0
0	1	1	0
1	0	0	0
1	0	1	0
1	1	0	0
1	1	1	1

This gives us the equation $\text{flag} = \overline{G_s} \oplus L_s C_{out}$

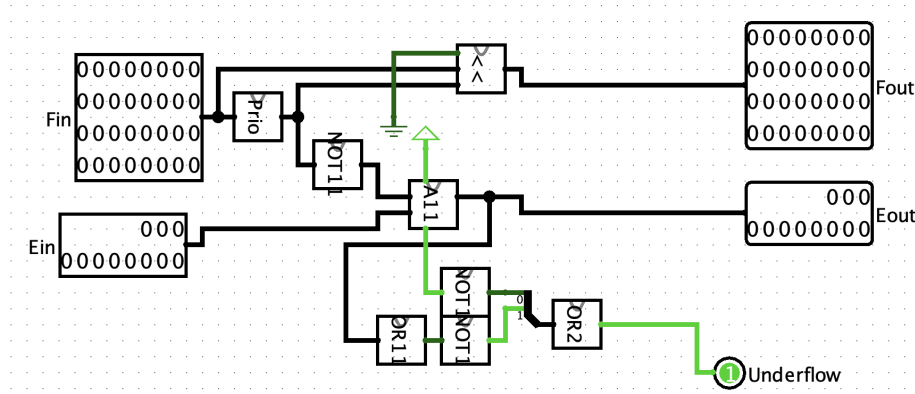


Figure 13: Normalizer

5.18 32-bit Zero

It gives an all-zero 32-bit output.




Figure 14: Zero32Bit

5.19 Rounder

This part rounds the mantissa. It determines the bit position in the 20-bit fraction where rounding needs to occur. Use a process, like a set of 3-bit truncators, to round the fraction based on the

identified position. If the bits beyond the rounding position are significant, add 1 to the rounded fraction. Adjust the exponent based on the rounding operation, considering any carry from the fraction rounding. Check for possible overflow due to rounding and adjust the exponent and fraction accordingly. Present the final result with the sign, adjusted exponent, and rounded fraction. The cases are as follows:

20 th bit	G	R	S	Action
X	0	X	X	Truncate
1	1	X	X	Round up
0	1	0	0	Truncate
X	1	1	X	Round up
X	1	X	1	Round up

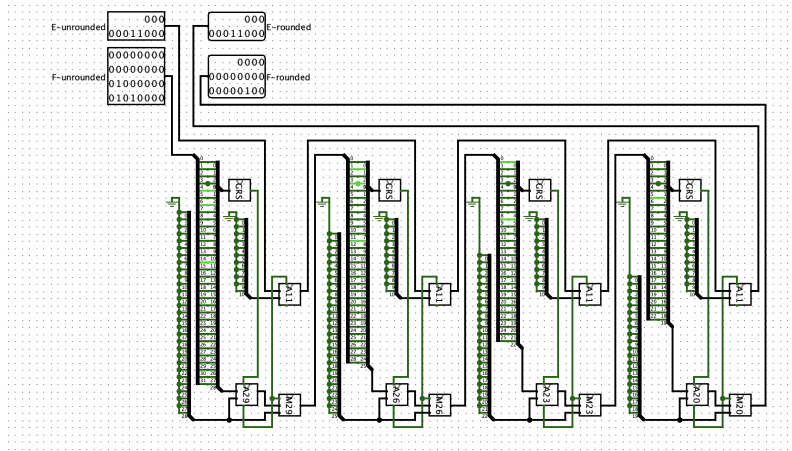


Figure 15: Rounder

5.20 Output Joiner

Finally, this circuit combines 3 incoming 1-bit, 11-bit, and 20-bit data into a single 32-bit output.




Figure 16: Output Joiner

5.21 Floating Point Adder

This is the full circuit of our floating point adder:

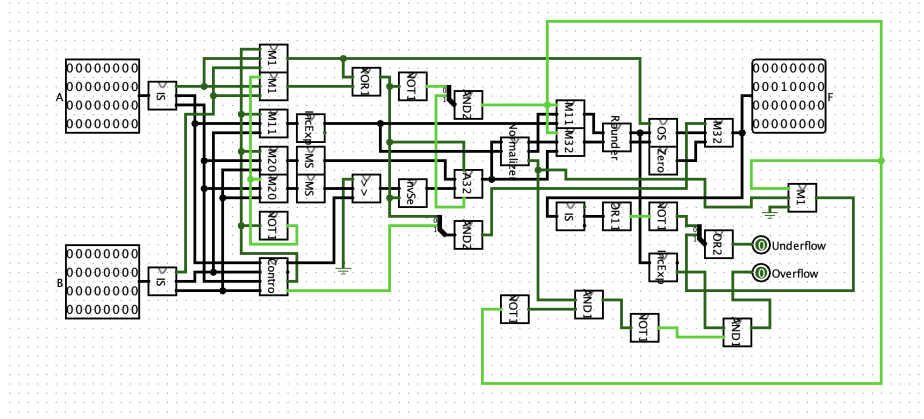


Figure 17: Floating Point Adder (FPA)




Figure 18: Complete Circuit

5.22 Final Circuit

5.23 Computing Sign Bit

When determining the sign bit, we took the sign bit of the larger (in value) number. We used the Control circuit output as the Mux selection bit to determine this sign bit.

5.24 Overflow and Underflow

Overflow and underflow flags were kept to signal that the output had exceeded the boundaries of normal numbers. For any two normal numbers, the highest allowable exponent is $2^{11} - 2$, hence the exponent of their sum can never exceed $2^{11} - 1$. Thus, checking if adding one to the final exponent results in a carry-out is enough to check for any overflow. On the other hand, for any two normal numbers, the lowest allowable exponent is 1, but since each of the fraction parts has 20 bits, the exponent of their difference can be equal to or lower than 0. Hence there are two cases: The exponent of the sum is zero, in that case checking if their NOR is 1 is sufficient; or the exponent of the sum has become lower than zero. In this case, we checked if the carry-out bit in the exponent decremter of the CustomLeftShifter is zero. If it is, this means that the sign of the

exponent has changed. Taking the OR of these two values determines whether the sum has caused an underflow.

5.25 Handling Zero or Small Input

If either of the inputs is zero, the other input will be passed directly to the output. Additionally, if the exponent of one of the inputs is significantly smaller than the other (a difference of more than 31), the input with the larger exponent is passed directly to the output.

6 ICs used with count as a chart

In this section, the integrated circuits (ICs) utilized in the implementation, along with their respective quantities, are presented in a concise chart:

IC Number	IC Name	Quantity
SN74HC08N	Quad 2-Input AND	90
SN74HC32N	Quad 2-Input OR	36
SN74HC86N	Quad 2-Input XOR	17
SN74HC04N	1-Input Hex-Inverter	68
SN74HC83N	4-bit Binary Full Adder	79
SN74HC157N	Quad 2 to 1 MUX	156

Table 2: ICs Used with Quantity

7 Simulator Used along with the Version Number

Software: Logisim

Version: Logisim-win-2.7.1

8 Discussion

For this assignment, our objective was to design a 32-bit FPA circuit capable of performing the addition of two floating point numbers operations. The software implementation presented challenges, particularly in strategically planning the placement of various modules.

- The entire setup was built using 7400-library integrated circuits, each serving specific functions.
- Our entire circuit is combinatorial, meaning it doesn't use any memory components like registers or flip-flops—just basic logic gates.
- To ensure both addends have the same exponent, we used a Custom Right Shifter. This device checks each bit of the shift amount value and decides whether to send the shifted or unshifted input value. For shift amounts greater than 31, we take the OR of the seven most significant bits, sending a full zero value if the result is 1.

- Inside the Normalizer, a Priority Encoder determines the position of the leftmost 1 in the 32-bit fraction. This value is used in CustomLeftShifter to shift the fraction by that amount and is subtracted from the exponent value.
- In the Rounder, we have cascaded 3-bit truncators, each truncating 3 bits from the fraction input. The exponent value is also rounded accordingly using adders.
- To avoid data loss and preserve even the lowest bit of information We considered the carry-out after addition as the leading bit of the result, allowing us to shift the sum leftwards without losing data.
- We handled addition or subtraction decisions using control, using the sign bit of the larger addend as the sign bit of the sum. For subtraction, the addend with the smaller absolute value is subtracted from the larger one.
- Rounding was meticulously done using cascading rounder devices.
- We put together individual components with smaller ones, although we didn't focus much on minimizing the number of ICs used.
- MUXs were extensively used for selecting and sorting values. Carry-out bits from different stages acted as selector bits for MUX in numerous instances.
- The circuit assumes that the addends will always be in normalized form, meaning no Zero, denormalized, NaN, or infinity values should be provided as inputs.